

Introduction: What is a Transaction

➡ Transactions:

- Units of work upon a data (persistent data, transient state)
- Bernstein: The execution of a program that performs an administrative function by accessing a shared database, usually on behalf of an on-line user.
- Examples
 - Reserve an airline seat and buy a ticket
 - Withdraw money from ATM
 - Verify a credit card sale
 - Carry out an order using on the Internet
 - Fire a missile
 - Download a video clip

What Makes Transactions Hard ?

- ➔ **Reliability** – system should rarely fail
- ➔ **Availability** – system must be up all the time
- ➔ **Response time** – system should not be slower
- ➔ **Throughput** – thousands of transactions per second
- ➔ **Scalability** – both small and Internet scale
- ➔ **Configurability** – above requirements and low cost
- ➔ **Atomicity** – no partial results
- ➔ **Durability** – a transaction is a legal contract
- ➔ **Distribution** – of users and data

What Makes Transactions Important ?

- ➔ **Well-defined and working abstraction at the level of databases.**
- ➔ **Most medium and large businesses use transactions for their production systems. The business cannot operate without it.**
- ➔ **A huge slice of the computer system market, over \$50B a year. Probably the single largest application of computers.**
- ➔ **Enables Internet commerce. In 2002, Intel did 10% or \$5B of contracting over electronic commerce.**

System Characteristics

➡ Typically < 100 transaction types per application

➡ Transaction size has high variance

- 0-30 disk accesses
- 10K - 1M instructions executed
- 2-20 messages

➡ A large scale example: airline reservations

- 150,000 active display devices
- thousands of disk drives
- 3000 transactions per second, peaks

Fault Tolerance: System Availability

➡ **Fault → Error → Failure**

➡ **Error latency**

➡ **Errors:**

➤ Latent

➤ Effective

➡ **Failures**

➤ Soft (recoverable)

➤ Hard

➡ **Mean time to failure (MTTF):**

➤ Average amount of time till the next failure

➡ **Mean time to repair (MTTR):**

➤ Average amount of time necessary for recovery

Fault Tolerance: System Availability

➔ Availability of the system:

$$\text{MTTF}/(\text{MTTF}+\text{MTTR})$$

➔ Some systems are *very* sensitive to downtime

➤ Airline reservation, stock exchange, telephone switching

<u>Downtime</u>	<u>Availability</u>
1 hour/day	95.8%
1 hour/week	99.41%
1 hour/month	99.86%
1 hour/year	99.9886%
1 hour/20years	99.99942%

Fault Tolerance: Error Avoidance and Correction

➡ Error correction

- Latent error processing
- Effective error processing
 - Failure correction or masking

➡ Avoidance: Valid construction

- ISO 9001, ...
- Duplexing, N-plexing (HW methods)
- Defensive programming
- Well-constructed software based on well-defined abstractions
 - ...
 - Transactions

Introduction: What Is a Transaction

➡ Transactions = ACID operations

➤ Atomicity:

- All or nothing property
- Includes all messages, operations, ...

➤ Consistency:

- A correct transformation of the state

➤ Isolation:

- A transaction is not aware of other transactions

➤ Durability:

- Once a transaction completes successfully (commits), its results survive some sort of failures

Atomicity

➡ All-or-nothing, no partial results

➤ Classical debit/credit example:

- A money transfer, debit one account, credit another
- Either debit and credit both run, or neither runs

➤ Successful completion is called **Commit**

➤ Transaction failure is called **Abort**

➡ Commit and abort are irrevocable actions (Durability)

➡ An Abort *undoes* operations that already executed

➤ For database operations, restore data's previous value from before the transaction

➤ But some real world operations are not undoable

- Examples: transfer money, print ticket, destroy house, fire missile
- Do such **real actions** at the end of transaction (as a part of commit)

Transactions: An Example

➔ Example 1: A very generic one

```
begin_work;  
    operation(object1);  
    operation(object2);  
    ...  
commit_work;
```

```
begin_work;  
    operation(object3);  
    operation(object4);  
    ...  
abort_work;
```

➔ Notes:

- Abort vs. Rollback
- Dependency on execution environment

Transactions: Another Example

➡ Example 2: SQL

```
SELECT id FROM accounts WHERE balance < 200;
```

```
INSERT INTO accounts VALUES ((67890, 1000));
```

```
UPDATE accounts SET balance=balance+1000 WHERE id=12345;
```

```
COMMIT_WORK;
```

Transactions: Yet Another Example

➡ Example 3: Enterprise JavaBeans (EJB)

```
try {  
    Account acc1 = homeInterface.create(...);  
    Account acc2 = homeInterface.create(...);  
    UserTransaction tx = JNDIService.lookup(...);  
    tx.begin();                // Thread becomes associated with tx  
    acc1.add(1000);            //  
    acc2.add(-1000);          // Operations performed on behalf of tx  
    tx.commit();  
} catch (TransactionAbortedException e) {  
    ...  
}
```

Example - ATM Dispenses Money a non-undoable operation

TX: Start

. . .

Transfer Money

Commit

System crashes
Transaction aborts
Money is transferred

TX: Start

. . .

Commit

Transfer Money

System crashes

*Deferred operation
never gets executed*

Reading Uncommitted Output Isn't Undoable

T1: Start

. . .

Display output

. . .

If error, Abort

▲ User reads output

...

▲ User enters input



Brain
transport

T2: Start

▲ **Get input from display**

. . .

Commit

Compensating Transactions

- ➡ **A transaction that reverses the effect of another transaction (which was committed)**
 - “Adjustment” in a financial system
 - Annul a marriage
- ➡ **Not all transactions have complete compensations**
 - Certain money transfers
 - Fire missile

Consistency

➡ Every transaction should maintain DB consistency

- Referential integrity - E.g. each order references an existing customer number and existing part numbers
- The books balance (debits = credits)

➡ ***Consistency preservation is a property of a transaction, not of the TP system*** (unlike the A, I, and D of ACID)

➡ If each transaction maintains consistency, then serial executions of transactions do too

- To prevent other transactions to see inconsistent temporary states, transactions are isolated
- Serializability theory

Isolation

➔ Concurrency control

- Intuitively, the effect of a set of transactions should be the same as if they ran independently
 - Formally, an interleaved execution of transactions is **serializable** if its effect is equivalent to a serial one
- Implies a user view where the system runs each user's transaction stand-alone
- Of course, transactions in fact run with lots of concurrency, to use device parallelism
- Serializability theory

➔ Locking-based schedulers

➔ Other schedulers

➔ Isolation levels

➔ Multiversion concurrency control

➔ ...

Durability

➔ Externalization of transaction's effects

- Storing modified data
- Sending messages
- Commit or Abort cannot be revoked

➔ Mostly implemented by resource managers

- Persistent stores
- Databases

Transaction Processing Monitors

- ➡ **A software product to create, execute and manage TP applications**
- ➡ **Takes an application written to process a single request and scales it up to a large, distributed system**
 - E.g. application developer writes programs to debit a checking account and verify a credit card purchase.
 - TP monitor helps system engineer deploy it to 10s/100s of servers and 10Ks of displays
- ➡ **Includes an application programming interface and tools for program development and system management**